

# Introduction to Machine Learning

Ning Xiong

Mälardalen University

# Agenda

- What is machine learning?
- Three classes of learning problems (supervised, unsupervised, reinforcement)
- Overview of some learning techniques
  - k-means clustering
  - linear regression
  - logistic regression
  - artificial neural network
  - Q-learning

# Part 1: What is Machine learning

# Inspiration from Human

Learn from experience



Learn from data



Follow instruction



# What is Machine Learning

Machine Learning is to **build computer programs** that can improve with experience at some task

- improve over some **Task T**
- with respect to a **performance measure P**
- based on **experience E**

# Applications of ML



- Learning to recognize spoken words
- Learning to drive autonomous vehicles (Google self-driving car)
- Learning to diagnose in medical environments
- Learning for diagnosis and prognosis for machine maintenance
- Learning for adaptive control of robots or industrial processes

# Three classes of learning problems

# Learning Problems

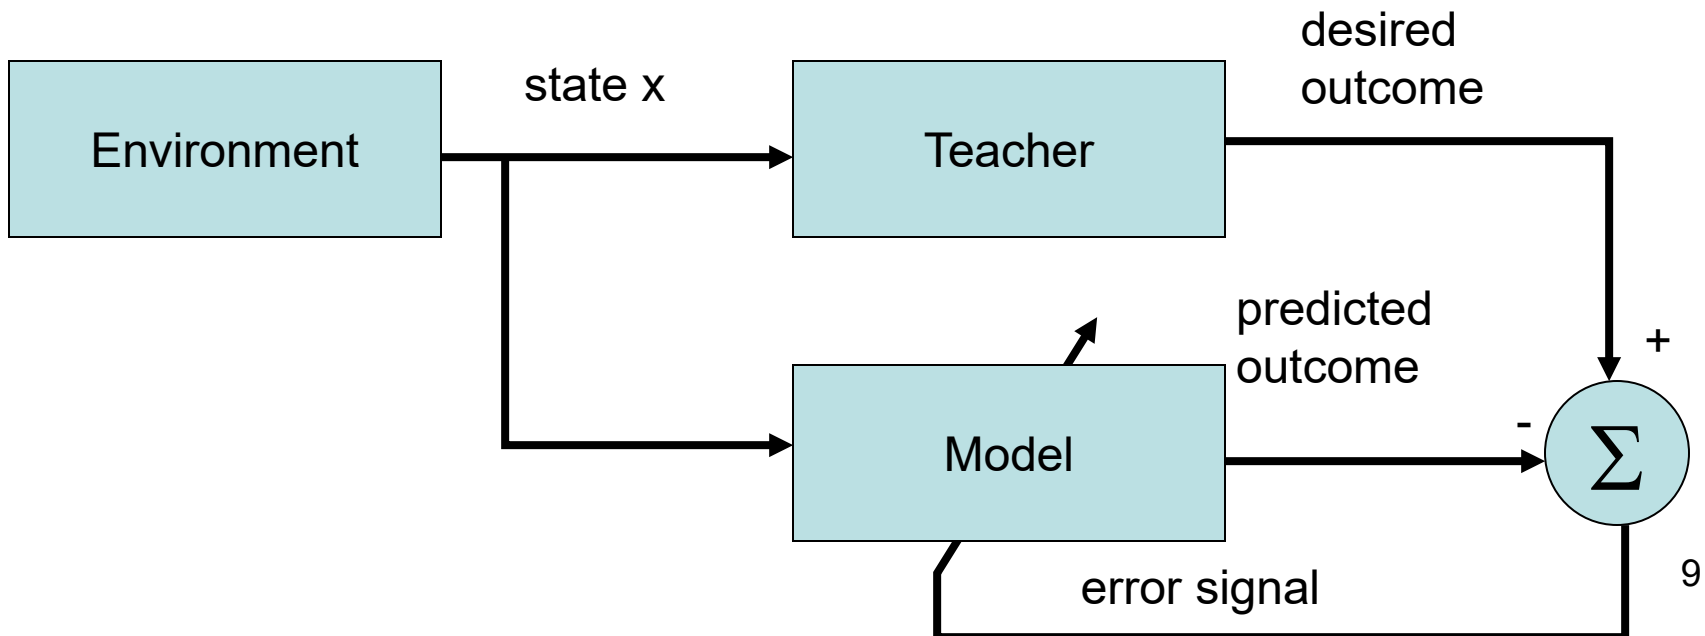
The nature of experience causes different kinds of learning problems

- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning



# Supervised Learning

- A teacher available to give desired outcome
- Direct experience represented by a set of input-output examples  $(\mathbf{x}_i, \mathbf{y}_i)$ , as training examples
- minimize the error between the predicted outcome of the learner and the desired outcome




# Unsupervised Learning

- No teacher or supervisor available
- No desired outputs in the experiences
- **Clustering**: Grouping similar instances
- **Data dimension reduction**: creating fewer new features

# Unsupervised Learning: Clustering


Objects similar in their attributes are clustered in the same group

Cluster 1



Outlier

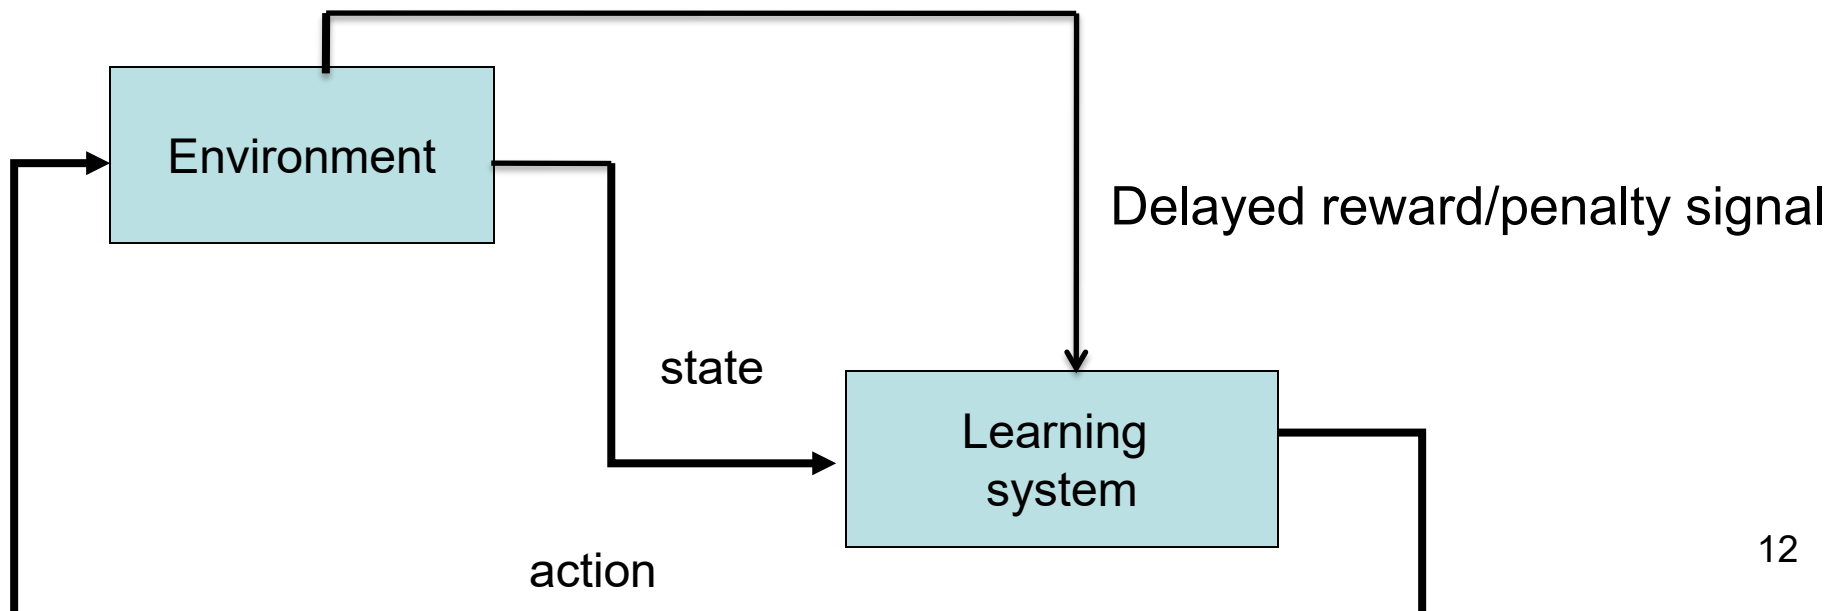
X



Cluster 2


# Reinforcement Learning

- Aim to obtain optimal action strategy without teacher
- learning through interaction with the environment
- exploration of states and actions
- goal: maximize accumulated rewards for the long term
- Indirect experience in form of delayed reward signal  
(temporal credit assignment problem)



# Credit Assignment Problem

Assigning credit or blame for the overall outcomes (delayed reward) to each of the internal decisions made by the learning machine which contributed to the overall outcomes



# Overview of learning techniques

# K-Means Clustering (unsupervised)


Task: divide the examples in to K clusters based on similarity

Initialize: number of clusters K; the initial centroids of the clusters


Iteration: assign data to clusters → recalculate the centroids




(a)




(b)




(c)



(d)




(e)



(f)

# Linear Regression (Supervised)

Build a linear function to model the relation of output variable  $y$  with the input variables  $x_1, \dots, x_n$



$$\hat{y} = w_0 + w_1 x_1 + \dots + w_n x_n$$

Finding the parameters to best fit the training samples

$$\min \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$



# Linear Regression

Suppose  $(x_{i1}, x_{i2}, \dots, x_{in}, y_i)$  ( $i=1 \dots m$ ) are the training examples.

Let

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \dots & x_{1n} \\ \dots & \dots & \dots & \dots \\ 1 & x_{m1} & \dots & x_{mn} \end{bmatrix} \quad \mathbf{Y} = \begin{bmatrix} y_1 \\ \dots \\ y_m \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} w_0 \\ \dots \\ w_n \end{bmatrix}$$

$$\begin{aligned} E(\mathbf{W}) &= \frac{1}{m} \|\mathbf{XW} - \mathbf{Y}\|^2 = \frac{1}{m} (\mathbf{XW} - \mathbf{Y})^T (\mathbf{XW} - \mathbf{Y}) \\ &= \frac{1}{m} (\mathbf{W}^T \mathbf{X}^T \mathbf{XW} - 2\mathbf{W}^T \mathbf{X}^T \mathbf{Y} + \mathbf{Y}^T \mathbf{Y}) \end{aligned}$$

$$\frac{\partial E}{\partial \mathbf{W}} = (2\mathbf{X}^T \mathbf{XW} - 2\mathbf{X}^T \mathbf{Y})/m = 0$$



$$\mathbf{W} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

Least square estimation

# Logistic Regression (Supervised)

- Build a linear surface to separate training examples into two classes
- Using linear surface to convert to class probability:  $z = \frac{e^{w_0 + w_1 x_1 + w_2 x_2}}{1 + e^{w_0 + w_1 x_1 + w_2 x_2}}$
- Finding parameters to minimize log loss function

$$E(W) = \begin{cases} -\log(Z) & \text{if } y = 1 \\ -\log(1 - z) & \text{if } y = 0 \end{cases}$$




# Gradient descent for learning

$$\text{Gradient} = \left[ \frac{\partial E}{\partial w_j} \right]$$

$$w_j = w_j - \eta \frac{\partial E}{\partial w_j}$$

$\eta$  is the learning rate specifying the step size in the gradient search



# Getting the gradient

Rewrite the loss function for a training example as

$$E(W) = -y \log(z) - (1 - y) \log(1 - z)$$

z: the predicted probability for the example

$$Z = \frac{e^{w_0 + \sum_{j=1}^n w_j x_j}}{1 + e^{w_0 + \sum_{j=1}^n w_j x_j}} = \frac{1}{1 + e^{-(w_0 + \sum_{j=1}^n w_j x_j)}}$$

y: the desired probability (1 or 0) for the example

$$\frac{\partial E}{\partial w_j} = (z - y)x_j$$

$$\frac{\partial E}{\partial w_0} = (z - y)$$

# Learning for logistic regression

Initialize parameters  $w_i$  with random values

Repeat



For each training example  $\{(x_1, \dots, x_n)^p, y^p\}$  do

1. Use the inputs  $(x_1, \dots, x_n)^p$  to calculate the predicted probability
2. Compute the gradients for all parameters
3. Update all the parameters based on gradients

until a certain criterion is satisfied

# Artificial Neural Network (ANN)


To directly represent complex nonlinear relation, we can build an artificial neural network consisting of many units.



- More details of ANN and its learning will be presented in Lecture 8


# Q learning (reinforcement)

- Goal: Learn optimal action policy  $\pi^*$ :  $s \rightarrow a$ , to maximize the total sum of rewards
- **Target function:**  $Q^*(s,a)$  referred as optimal value of action in state  $S$



- $Q^*(s, a)$  is defined as the accumulated reward that will be obtained from state  $s$  by first doing action  $a$  then following optimal policy
- Optimal decision making:  $\pi^*(s) = \arg \max_{a \in A(s)} Q^*(s, a)$

# Q learning rule



$$Q^*(s, a) = r + \gamma \max_{\forall a'} Q^*(s', a')$$



Q learning rule:  $Q(s, a) = r + \gamma \max_{\forall a'} Q(s', a')$

Through many interactions:  $Q(s, a) \rightarrow Q^*(s, a)$



# A table for estimates of $Q^*$ values

	$a_1$	$a_2$	....	$a_m$
$s_1$	*	*	*	*
$s_2$	*	*	*	*
...	*	*	*	*
$s_n$	*	*	*	*

- Each time update one entity in the table following the Q-learning rule